

Shishir Lohar

U.S. Citizen | (484)-784-8419 | shishirlohar@gmail.com | linkedin.com/in/shishirlohar | [Portfolio](#) | [GitHub](#)

EDUCATION

University of Georgia, Franklin College of Arts and Sciences

December 2026

B.S. in Computer Science, Focus in Artificial Intelligence; GPA: 3.5/4.0

Athens, GA

Relevant Coursework:

Software Engineering, Data Structures, Algorithms, Systems-Level Programming, Web Programming, Robotics

TECHNICAL SKILLS

Languages: Python, Java, C/C++, JavaScript, TypeScript, SQL, HTML/CSS

Frameworks/Libraries: React, Next.js, Node.js, FastAPI, Flask, Spring Boot, Streamlit, Leaflet/Folium

AI/Data: OpenAI API, MCP, Llama, Qwen, scikit-learn, Pandas

Databases & Integrations: PostgreSQL, SQLite, REST APIs, Discord Webhooks

Tools/Cloud: Git, Docker, Railway, Vercel, AWS (S3, CloudFront), GCP, GitHub Actions, Kafka, Postman, Linear

EXPERIENCE

Lead Software Engineer | Kappa Theta Pi | Athens, GA

September 2026 – Present

- Leading a 3-person engineering team developing KTPI, a deployment investigation platform using Python, FastAPI, and PostgreSQL to ingest GitHub workflow and deployment events through REST APIs and display incident history and reports in a React dashboard
- Developing an agentic investigator with the OpenAI API and MCP tools to analyze workflow logs, commits, and test failures, generating evidence-backed root-cause hypotheses and recommended fixes
- Designing Docker services on our local server and React hosting on AWS S3/CloudFront, with GitHub Actions CI/CD, Discord notifications, event retries and deduplication, automated tests, and peer code reviews

Software Engineer Intern | UGA Warnell School of Forestry | Athens, GA

April 2026 – Present

- Engineered an AI assistant using Python and local LLMs to convert natural-language questions and prompts into validated SQL database queries and interactive map visualizations across 100+ test prompts
- Processed 500K+ rows across large forestry datasets using Pandas, combining data validation, filtering, and the AI model Qwen2.5's output to create visualizations using Streamlit in under 30 seconds for workflows that could take up to 1 hour to complete manually using traditional data tools, and documented updates with Linear
- Evaluated 4 Llama and Qwen models and deployed with Docker on Railway to run the AI assistant, integrating persistent storage, deterministic fallback logic, and visualization with Leaflet/Folium following CI/CD

Founder & Software Engineer | Spotr | Duluth, GA

May – August 2026

- Secured \$7500+ in funding to build an AI-powered vehicle sourcing platform for dealerships, using 3+ automotive data APIs to search tens of thousands of listings across hundreds of dealerships and online marketplaces
- Developed API-driven workflows to retrieve, process, filter, and rank vehicle inventory across 10+ dealership-defined criteria, consolidating multiple external data sources into a unified search experience

PROJECTS

Software Engineer (Forage Experience) | JPMorgan Chase & Co. | Remote

February – March 2026

- Engineered a Spring Boot service to receive, validate, and record high-volume financial transactions, integrating Kafka-based processing with a SQL database for persistent storage
- Developed a REST API to expose processed financial data, enabling reliable downstream access and communication with external service consumers

Machine-Learning Health Model | [GitHub Repo](#)

August – September 2025

Python, scikit-learn, Pandas, Streamlit, Random Forest

- Engineered an end-to-end health trend prediction pipeline across 1,000 synthetic health records and 8 biometric and lifestyle features, including sleep, activity, caloric intake, blood pressure, and resting heart rate
- Trained and evaluated a 100-tree Random Forest classifier with scikit-learn and Pandas using an 80/20 train-test split, achieving 98% classification accuracy across improving, stable, and declining health-trend categories
- Serialized the trained model with Joblib and integrated real-time inference into an interactive Streamlit web application, enabling users to enter daily health metrics and receive immediate personalized trend predictions